

Help Website Q&A Agent

GitHub - <https://github.com/reyanalam/Al-powered-question-answering-agent.git> (Private , only access to vatsal@pulsegen.io)

Video -

<https://drive.google.com/file/d/1GodG3DrcDPRtFCCYmCaFbvQjxz6cMYeo/view?usp=sharing>

This project is an AI-powered question-answering system that processes help website documentation and answers user queries related to product features, integrations, and functionalities.

Technical Architecture Overview

The system architecture follows a modular pipeline, ensuring clarity, extensibility, and ease of debugging. The major components include:

1. Web Crawler

- **Purpose:** Extracts meaningful and structured data from static HTML content on help websites.
- **Technologies Used:**
 - `requests` – for HTTP requests
 - `beautifulsoup4` – for parsing and extracting HTML content
 - `urllib3`, `tldextract` – for robust URL handling
- **Features:**
 - Recursive crawling with customizable depth
 - Filters out irrelevant links and external domains
 - Collects paragraphs, lists, tables, and headers

2. Preprocessing Module

- **Purpose:** Cleans and tokenizes the content for further semantic analysis.
- **Technology Used:**
 - spaCy with the large English model
- **Steps:**
 - Stop word removal
 - Lemmatization

3. Chunking Engine

- **Purpose:** Segments text into semantically similar chunks.
- **Approach:**
 - Uses cosine similarity between token vectors
 - Groups tokens until similarity drops below a defined threshold
 - Ensures each chunk is meaningful and self-contained

4. Semantic Search

- **Purpose:** Embeds user queries and document chunks into vector space for comparison.
- **Approach:**
 - Averages spaCy token vectors to represent full queries and chunks
 - Computes cosine similarity between the query vector and each chunk
 - Returns top 3 most relevant chunks

Implementation Approach

1. **Input Handling:** User provides the base URL of a help website (e.g., <https://help.zluri.com/>).
 2. **Crawling:** Internal pages are crawled recursively up to a user-defined depth.
 3. **Extraction & Cleaning:** HTML content is parsed and non-relevant tags are discarded.
 4. **Text Preprocessing:** The clean text is tokenized, lemmatized using spaCy.
 5. **Chunking:** Text is chunked based on semantic cohesion measured by cosine similarity.
 6. **Vector Embedding & Search:** Both chunks and queries are embedded and compared using cosine similarity.
 7. **Answer Retrieval:** Top 3 matching chunks are returned as answers to the user query.
-

Testing Approach

- **Unit Testing:**
 - Core modules like web crawler, chunker, and vector generator are tested individually.
 - `test.py` script verifies overall functionality using a sample URL.
- **Manual Verification:**
 - Compare returned answers with original documentation content to verify relevance and accuracy.

Future Improvement Suggestions

1. **Enhanced Chunking**
 - Replace custom chunking with LangChain or other NLP-based segmenters for improved semantic segmentation.
2. **Transformer-based Embeddings**

- Use models like BERT, RoBERTa, or SentenceTransformers for contextual embeddings to boost answer quality.

3. Improved Search Algorithms

- Implement FAISS or Elasticsearch for scalable and faster vector-based search over large corpora.

4. Support for Dynamic Content

- Integrate Selenium or Playwright to scrape JavaScript-heavy websites.

5. Web Interface

- Add a frontend UI for better user interaction instead of terminal-based input.